

SCSE3040 Machine Learning Operations --- Practical 03 submission

Name	Shaurya Nigam
Roll number	S24CSEU0497
Practical	P03 --- Choosing a Model Honestly
Notebook	P03_S24CSEU0497.ipynb
Self-check	8 / 8 PASS

Practical 03 --- Choosing a Model Honestly

****Overfitting, the train/test gap, and cross-validation****

SCSE3040 Machine Learning Operations &middledot; Bennett University
&middledot; Session 2026-27

```
| | |  
|---|---|  
| **Lectures this follows** | L04-L05 |  
| **Course Outcome** | C02 |  
| **Lab duration** | 120 minutes |  
| **Memory needed** | about 400 MB (fine on a 4 GB or 8 GB machine)  
|  
| **Extra software** | nothing beyond the course venv |  
| **Marks** | 10 |
```

Aim

1. Train three different models on the same data and compare them fairly.
 2. See overfitting happen, and measure it as a gap between two scores.
 3. Use cross-validation instead of trusting one lucky split.
 4. Choose a model for a reason you can defend, not because it scored best.
-

What you need to know first

In P02 you trained one model and it did well. A reasonable next thought is "let me try a more powerful model and do even better". Today you find out why that thought is dangerous.

A **decision tree** is a model that asks yes/no questions: *is the distance more than 6 km? is it raining?* Ask enough questions and you can separate every single training order from every other one. At that point the model has not learned a pattern --- it has memorised the training data. It will look perfect in training and fall apart on anything new.

That failure has a name: **overfitting**. And it has a measurement. Score the model twice, once on the data it trained on and once on data it has never seen. The distance between those two numbers is the **train/test gap**. A small gap means the model learned something general. A big gap means it memorised.

There is one more problem to fix. In P02 your test set was one particular 120 orders, chosen by one particular seed. What if those 120 happened to be easy? **Cross-validation** removes that luck: it cuts the data into 5 equal parts, trains 5 times, and each time tests on a different

part. You get 5 scores instead of 1, and you report the average. It costs 5 times the computing time, which on this dataset means well under a second.

The habit to build today: *the best score is not the best model*. A simple model that scores almost as well is usually the one to ship, because you can explain it, it trains faster, and it breaks less.

Before you start

- Practical P02 is finished. You can split data, train a model and read an MAE.

- You remember that the test set is the only honest score.

How to run a cell

Click on a grey code cell, then press **Shift + Enter**. The cell runs and the cursor moves to the next one. A number appears in the `[ ]` on the left when the cell has finished.

Run the cells in order, from the top. A later cell almost always uses something an earlier cell created. If you jump ahead you will see a `NameError`, which just means "you have not made that thing yet".

If everything goes wrong, use the menu: **Kernel -> Restart Kernel and Clear All Outputs**, then start again from the first cell. Nothing is damaged by doing this.

```
In [1]: # Step 0 --- check the workbench before we start.
# This cell only looks; it changes nothing. Run it and read the last line.

import sys
from pathlib import Path

print("Python :", sys.version.split()[0])
print("Folder :", Path.cwd().name)

_missing = []
for _name in ['numpy', 'pandas', 'sklearn', 'matplotlib']:
    try:
        __import__(_name)
    except ImportError:
        _missing.append(_name)

for _name in ['numpy', 'pandas', 'sklearn', 'matplotlib']:
    _mark = "missing" if _name in _missing else "ok"
    print(f" {_name:<14} {_mark}")

if _missing:
    print()
    print("STOP. Some libraries are missing:", ", ".join(_missing))
    print("Ask your instructor to run the setup in labs/SETUP.md.")
else:
    print()
    print("All good. You can carry on to Step 1.")
```

```
Python : 3.11.9
Folder  : P03-model-choice
numpy   ok
pandas  ok
sklearn ok
matplotlib ok
```

All good. You can carry on to Step 1.

Step 0b --- the dataset

Every practical in this course uses the same 600 food deliveries. The next cell makes sure the file is there.

```

In [2]: # The delivery-time dataset every practical in this course uses.
# If the file is missing we build it again from the same seed, so every
# student in the room gets byte-for-byte the same 600 rows.

import csv
from pathlib import Path

import numpy as np

SEED = 42
N_ROWS = 600
DATA = Path("../") / "data" / "delivery_times.csv"

def make_delivery_csv(path=DATA):
    """Write the 600-row delivery dataset. Same formula as the lectures."""
    rng = np.random.default_rng(SEED)
    distance_km = np.round(rng.uniform(0.5, 12.0, N_ROWS), 2)
    prep_time_min = np.round(rng.uniform(5, 30, N_ROWS), 0)
    traffic_level = rng.integers(1, 4, N_ROWS)
    rain = rng.binomial(1, 0.25, N_ROWS)
    delivery_min = np.round(
        6.0
        + 3.1 * distance_km
        + 0.65 * prep_time_min
        + 4.2 * traffic_level
        + 5.5 * rain
        + rng.normal(0, 2.5, N_ROWS),
        1,
    )
    path.parent.mkdir(parents=True, exist_ok=True)
    with path.open("w", newline="", encoding="utf-8") as fh:
        w = csv.writer(fh)
        w.writerow(["distance_km", "prep_time_min", "traffic_level",
                    "rain", "delivery_min"])
        for i in range(N_ROWS):
            w.writerow([distance_km[i], int(prep_time_min[i]),
                        int(traffic_level[i]), int(rain[i]), delivery_min[i]])
    return path

if not DATA.exists():
    make_delivery_csv()
    print("dataset rebuilt ->", DATA)
else:
    print("dataset found ->", DATA)

```

```
dataset found -> ../data\delivery_times.csv
```

Walkthrough

Read each step, then run its cell.

Step 1 --- Set up the data once, for all three models

Same data, same split, same seed as P02. Every model today gets exactly the same deal --- otherwise the comparison is worthless.

```
In [3]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split

orders = pd.read_csv(DATA)

FEATURES = ["distance_km", "prep_time_min", "traffic_level", "rain"]
X = orders[FEATURES]
y = orders["delivery_min"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

print("train:", len(X_train), " test:", len(X_test))
```

train: 480 test: 120

Step 2 --- A helper that scores a model both ways

We will do the same two measurements many times today, so write them once. This function returns the training MAE, the test MAE, and the gap between them.

```
In [4]: from sklearn.metrics import mean_absolute_error

def score_both_ways(model, name):
    """Train the model, then score it on train and on test."""
    model.fit(X_train, y_train)
    train_mae = mean_absolute_error(y_train, model.predict(X_train))
    test_mae = mean_absolute_error(y_test, model.predict(X_test))
    gap = test_mae - train_mae
    print(f"{name:<26} train {train_mae:5.2f} "
          f"test {test_mae:5.2f} gap {gap:5.2f}")
    return {"name": name, "train": train_mae,
            "test": test_mae, "gap": gap}

print("helper ready")
```

helper ready

Step 3 --- Model 1: linear regression, the honest baseline

The model from P02. Watch the gap: it should be tiny, because a straight line is too simple to memorise anything.

```
In [5]: from sklearn.linear_model import LinearRegression

linear = score_both_ways(LinearRegression(), "LinearRegression")
```

LinearRegression train 2.04 test 1.92 gap -0.11

Train and test scores almost equal. This model is not overfitting. It may not be the most accurate thing possible, but it is telling you the truth.

Step 4 --- Model 2: a decision tree with no limits

Now the dangerous one. A decision tree with no maximum depth is allowed to keep asking questions until every training order sits in its own little box.

Look at the training score first. Then look at the test score.

```
In [6]: from sklearn.tree import DecisionTreeRegressor

wild_tree = score_both_ways(
    DecisionTreeRegressor(random_state=42), "DecisionTree (no limit)"
)
```

DecisionTree (no limit) train 0.00 test 3.43 gap 3.43

A training MAE near **zero** and a much worse test MAE. The model got every training order perfectly right and learned nothing useful.

That gap is overfitting. You have now seen it happen, not just been told about it.

Step 5 --- Model 3: the same tree, but shallow

Now stop the tree early. `max_depth=4` means it may only ask four questions before it has to answer. It can no longer memorise.

```
In [7]: small_tree = score_both_ways(
    DecisionTreeRegressor(max_depth=4, random_state=42),
    "DecisionTree (depth 4)",
)
```

DecisionTree (depth 4) train 3.66 test 4.23 gap 0.57

Worse on training, better on test, and a far smaller gap. Limiting a model's freedom made it *more* useful. That idea is the whole of model selection.

Step 6 --- Model 4: a random forest

A **random forest** grows many different trees and averages their answers. Individual trees still overfit, but their mistakes point in different directions and largely cancel out.

We use 50 trees. More would be slower and no better here.

```
In [8]: from sklearn.ensemble import RandomForestRegressor

forest = score_both_ways(
    RandomForestRegressor(n_estimators=50, random_state=42),
    "RandomForest (50 trees)",
)
```

RandomForest (50 trees) train 1.00 test 2.39 gap 1.39

Step 7 --- Put the four results side by side

Never compare models from memory. Put them in a table.

```
In [9]: results = pd.DataFrame([linear, wild_tree, small_tree, forest])
results = results.round(2).sort_values("test")

print(results.to_string(index=False))
print()
best = results.iloc[0]
print(f"Lowest test MAE: {best['name']} at {best['test']:.2f} minutes")
```

	name	train	test	gap
	LinearRegression	2.04	1.92	-0.11
	RandomForest (50 trees)	1.00	2.39	1.39
	DecisionTree (no limit)	0.00	3.43	3.43
	DecisionTree (depth 4)	3.66	4.23	0.57

Lowest test MAE: LinearRegression at 1.92 minutes

Read the `gap` column as carefully as the `test` column. A model with a good test score and a huge gap got lucky, and luck does not survive contact with real users.

Step 8 --- Draw the gap, because a table makes you read


The table is correct but slow --- you must read eight numbers and subtract in your head. A chart does the subtraction for your eyes.

matplotlib is Python's standard charting library. It has been sitting in your course environment since P01; today it earns its place. The first line, `%matplotlib inline`, tells Jupyter to draw the picture right below the cell.


```
In [10]: %matplotlib inline
import matplotlib.pyplot as plt

four = [linear, wild_tree, small_tree, forest]
xs = np.arange(len(four))

fig, ax = plt.subplots(figsize=(9, 3.8))
ax.bar(xs - 0.2, [m["train"] for m in four], 0.38,
       label="train MAE", color="#6B7785")
ax.bar(xs + 0.2, [m["test"] for m in four], 0.38,
       label="test MAE", color="#9E3B33")
ax.set_xticks(xs, [m["name"] for m in four], fontsize=9)
ax.set_ylabel("MAE (minutes)")
ax.set_title("The distance between a pair of bars is the gap")
ax.legend()
plt.tight_layout()
plt.show()
```



Look at the pair for the no-limit tree. Its grey train bar has **vanished** --- an MAE of 0.00 draws as no bar at all --- while its red test bar is among the tallest. Bars far apart mean the model memorised; bars close together mean it learned.

The linear model's two bars are the closest pair on the chart. That is what honesty looks like.

Step 9 --- Stop trusting one split: cross-validation

Every number above came from one 120-order test set. Change the seed and they all move a little.

5-fold cross-validation cuts the data into 5 parts and trains 5 times, testing on a different part each time. We report the average and the spread.

scikit-learn reports errors as *negative* numbers, because its convention is that bigger is better. We flip the sign back.

```
In [11]: from sklearn.model_selection import cross_val_score

def cross_validate(model, name):
    """Average MAE over 5 different train/test splits."""
    scores = -cross_val_score(
        model, X, y, cv=5, scoring="neg_mean_absolute_error"
    )
    print(f"{name:<26} MAE {scores.mean():5.2f} "
          f"+/- {scores.std():4.2f}")
    return scores.mean()

cv_linear = cross_validate(LinearRegression(), "LinearRegression")
cv_tree = cross_validate(
    DecisionTreeRegressor(max_depth=4, random_state=42),
    "DecisionTree (depth 4)")
cv_forest = cross_validate(
    RandomForestRegressor(n_estimators=50, random_state=42),
    "RandomForest (50 trees)")
```

LinearRegression	MAE	2.03 (+/- 0.16)
DecisionTree (depth 4)	MAE	4.57 (+/- 0.18)
RandomForest (50 trees)	MAE	2.69 (+/- 0.40)

The +/- is the spread across the 5 folds. If two models are closer together than their spread, you cannot honestly say one is better.

Step 10 --- Tune one setting, properly

How deep should the tree be? Do not guess, and do not try depths on the test set --- that would leak the answers.

Try each depth with cross-validation, which never touches the test set at all.

```
In [12]: depths = [2, 3, 4, 6, 8, None]
scores_by_depth = {}

for depth in depths:
    tree = DecisionTreeRegressor(max_depth=depth, random_state=42)
    mae = -cross_val_score(
        tree, X, y, cv=5, scoring="neg_mean_absolute_error"
    ).mean()
    scores_by_depth[depth] = round(mae, 3)
    print(f"max_depth={str(depth):<5} MAE {mae:5.2f}")

best_depth = min(scores_by_depth, key=scores_by_depth.get)
print()
print("best depth:", best_depth)
```

max_depth=2	MAE	5.86
max_depth=3	MAE	4.85
max_depth=4	MAE	4.57
max_depth=6	MAE	3.64
max_depth=8	MAE	3.39
max_depth=None	MAE	3.48

best depth: 8

Read the **shape**, not just the winner.

Depth 2 is the worst of all: too simple to capture anything, which is called **underfitting**. The scores improve as the tree is allowed more room, reach their best in the middle, and then get slightly worse again when the limit is removed altogether.

Too rigid and too free are both bad. The good answer is in between, and the only way to find it is to measure --- which is what you just did, without ever touching the test set.


Step 11 --- Draw the sweep, because ten numbers hide a shape

The loop printed six scores. Their *shape* is the real lesson, and a shape is easier to see than to read. Draw MAE against depth.

The red circle marks the depth cross-validation liked best.

```
In [13]: xs = np.arange(len(depths))
maes = [scores_by_depth[d] for d in depths]
best_i = int(np.argmin(maes))

fig, ax = plt.subplots(figsize=(8, 3.6))
ax.plot(xs, maes, marker="o", color="#2C5F79", linewidth=2)
ax.plot(xs[best_i], maes[best_i], marker="o", markersize=13,
        markerfacecolor="none", markeredgewidth=2,
        markeredgecolor="#9E3B33")
ax.annotate(f"best: depth {depths[best_i]}",
            (xs[best_i], maes[best_i]),
            xytext=(0, -22), textcoords="offset points",
            ha="center", color="#9E3B33")
ax.set_xticks(xs, [str(d) for d in depths])
ax.set_ylim(bottom=min(maes) - 0.8)
ax.set_xlabel("max_depth")
ax.set_ylabel("cross-validated MAE (minutes)")
ax.set_title("Too rigid on the left, too free on the right")
plt.tight_layout()
plt.show()
```



Now the shape is unmistakable. The left end towers over everything: depth 2 is the worst score on the whole chart --- **underfitting**, a model too rigid to learn anything. The curve falls as the tree earns more freedom, bottoms out in the middle, then ticks back up when the last limit is removed.

Both ends of the curve are bad. The answer lives at the bottom, and you found it without ever touching the test set.

Step 12 --- Choose, and be able to say why

Now make the decision the way a professional would, out loud.

```
In [14]: print("Cross-validated MAE, best first:")
for name, mae in sorted(
    {"LinearRegression": cv_linear,
     "DecisionTree(4)": cv_tree,
     "RandomForest(50)": cv_forest}.items(),
    key=lambda kv: kv[1],
):
    print(f" {name:<20} {mae:5.2f}")

print()
print("Our choice for the delivery service: LinearRegression.")
print("Reasons:")
print("  1. It has the lowest cross-validated error here.")
print("  2. It trains in milliseconds and needs almost no memory.")
print("  3. We can read its coefficients and explain a prediction")
print("      to a restaurant owner. Nobody can read 50 trees.")
print()
print("Be careful with reason 1. This dataset really was built")
print("from a straight line, so a straight line wins. On messier")
print("real data the forest usually wins instead -- and then you")
print("must decide what its extra accuracy is worth, against")
print("reasons 2 and 3, which do not change.")
```

Cross-validated MAE, best first:

LinearRegression	2.03
RandomForest(50)	2.69
DecisionTree(4)	4.57

Our choice for the delivery service: LinearRegression.

Reasons:

1. It has the lowest cross-validated error here.
2. It trains in milliseconds and needs almost no memory.
3. We can read its coefficients and explain a prediction to a restaurant owner. Nobody can read 50 trees.

Be careful with reason 1. This dataset really was built from a straight line, so a straight line wins. On messier real data the forest usually wins instead -- and then you must decide what its extra accuracy is worth, against reasons 2 and 3, which do not change.

You will not always choose the simple model. But you must always be able to say why you chose what you chose. "It scored best" is not a reason a senior engineer will accept on its own --- and, as the last paragraph says, it will not always be true either.

Your turn

The walkthrough above is finished. Now you write some code.

There are **3 tasks**. Each one is small. Each one has a hint. Do them in order.

Where you see `# TODO`, replace that line with your own code. Do not delete the variable name on the left of the `=` sign --- the self-check at the bottom looks for exactly that name.

When you have tried all three, run the **self-check** cell at the end. It prints a table telling you which tasks are correct. You can run it as many times as you like.

Task T1 --- Measure overfitting yourself

Train a `DecisionTreeRegressor` with `max_depth=12` and `random_state=42`, and work out its train/test gap by hand.

Store the three numbers in `T1_train_mae`, `T1_test_mae` and `T1_gap`.

`T1_gap` must be the test MAE **minus** the train MAE.

Hint. You may not use `score_both_ways` for this one -- write the three lines out, so you can see what the helper was doing. Fit the tree, then call `mean_absolute_error` twice: once with `y_train` and `X_train`, once with `y_test` and `X_test`.

```
In [15]: deep_tree = DecisionTreeRegressor(max_depth=12, random_state=42)

# Train it on the training orders only. The test set stays untouched.
deep_tree.fit(X_train, y_train)

# Score it twice, on the very same model.
# First on the 480 orders it has already seen while learning.
T1_train_mae = mean_absolute_error(y_train, deep_tree.predict(X_train))

# Then on the 120 orders it has never seen. This is the honest score.
T1_test_mae = mean_absolute_error(y_test, deep_tree.predict(X_test))

# The gap is how much worse it does on new orders than on old ones.
T1_gap = T1_test_mae - T1_train_mae

print(f"train {T1_train_mae:.2f} test {T1_test_mae:.2f} gap {T1_gap:.2f}")
```

```
train 0.05 test 3.40 gap 3.35
```

Task T2 --- Tune the forest, without touching the test set

A random forest has a setting `max_depth` too. Use ****5-fold cross-validation**** to find which of `[2, 4, 6, None]` is best for a `RandomForestRegressor(n_estimators=50, random_state=42)`.

Build a dictionary `T2_scores` mapping each depth to its cross-validated MAE, then put the winning depth in `T2_best_depth`.

You must not use `X_test` or `y_test` anywhere in this task.

Hint. Step 10 does exactly this for a single tree. Copy that loop and change two things: the model, and the list of depths. The winner is `min(T2_scores, key=T2_scores.get)`.

```
In [16]: T2_depths = [2, 4, 6, None]

# Cross-validate a 50-tree forest at each depth, exactly like Step 10 did
# for a single tree. cross_val_score returns negative MAE because its
# convention is "bigger is better", so we put the minus sign back.
T2_scores = {}

for depth in T2_depths:
    forest_at_depth = RandomForestRegressor(
        n_estimators=50, max_depth=depth, random_state=42
    )
    mae = -cross_val_score(
        forest_at_depth, X, y, cv=5, scoring="neg_mean_absolute_error"
    ).mean()
    T2_scores[depth] = round(float(mae), 3)
    print(f"max_depth={str(depth):<5} MAE {mae:5.2f}")

# Lowest MAE wins. Note X_test and y_test appear nowhere in this cell --
# the 5 folds are cut out of X and y, so nothing leaked.
T2_best_depth = min(T2_scores, key=T2_scores.get)

print()
print(T2_scores)
print("best:", T2_best_depth)
```

```
max_depth=2      MAE   5.27
max_depth=4      MAE   3.64
max_depth=6      MAE   2.84
max_depth=None   MAE   2.69
```

```
{2: 5.271, 4: 3.643, 6: 2.844, None: 2.688}
best: None
```

Task T3 --- A comparison you can reuse

Write a function `compare(models)` that takes a **dictionary** of `{name: model}` and returns a dictionary of `{name: test MAE}`.

It must fit each model on `X_train`/`y_train` and score it on `X_test`/`y_test`.

Then call it on the three models given in the stub and store the answer in `T3_table`.

Hint. Loop over `models.items()`, call `.fit(...)`, then `mean_absolute_error(y_test, m.predict(X_test))`, and collect the results into a new dictionary that you `return` at the end.

```
In [17]: def compare(models):
          """Fit every model on the same training data, score it on the same
          test set, and hand back a {name: test MAE} dictionary."""
          table = {}
          for name, model in models.items():
              model.fit(X_train, y_train)
              table[name] = mean_absolute_error(y_test, model.predict(X_test))
          return table

          T3_table = compare({
              "linear": LinearRegression(),
              "tree4": DecisionTreeRegressor(max_depth=4, random_state=42),
              "forest": RandomForestRegressor(n_estimators=50, random_state=42),
          })

          # Print it best-first so the comparison reads at a glance.
          for name, mae in sorted(T3_table.items(), key=lambda kv: kv[1]):
              print(f" {name:<8} test MAE {mae:5.2f} minutes")
```

```
linear    test MAE   1.92 minutes
forest    test MAE   2.39 minutes
tree4     test MAE   4.23 minutes
```

Self-check

Run the cell below to mark your work.

```

In [18]: # -----
# SELF-CHECK --- run this when you have attempted the tasks above.
# It never breaks your notebook. A task you have not done yet simply
# shows FAIL.
# -----

_results = []

def _check(label, fn):
    """Evaluate one graded condition without ever raising."""
    try:
        ok = bool(fn())
    except Exception:
        ok = False
    _results.append((label, ok))

_check('T1 | the depth-12 tree is trained', lambda: hasattr(deep_tree, 'tree_'))
_check('T1 | T1_train_mae and T1_test_mae are that tree's two scores', lambda: abs(T1_train_mae - T1_test_mae) < 0.001)
_check('T1 | T1_gap is test minus train, and it is positive', lambda: abs(float(T1_test_mae - T1_train_mae)) < 0.001)
_check('T2 | T2_scores has one score per depth', lambda: set(T2_scores) == {2, 4, 6, 8, 10, 12})
_check('T2 | every score is a plausible MAE in minutes', lambda: all(0 < float(v) < 120 for v in T2_scores.values()))
_check('T2 | T2_best_depth is the depth with the lowest score', lambda: T2_best_depth == min(T2_scores.keys(), key=lambda d: T2_scores[d]))
_check('T3 | compare() returns one score per model', lambda: set(T3_table) == {'linear', 'walkthrough', 'deep_tree'})
_check('T3 | the linear score matches the walkthrough', lambda: abs(float(T3_table['linear']) - float(T3_table['walkthrough'])) < 0.001)

print("=====")
print("SELF-CHECK   Practical 03 --- Choosing a Model Honestly")
print("=====")
for _label, _ok in _results:
    print(f"  [{'PASS' if _ok else 'FAIL'}]  {_label}")
print("-----")
_passed = sum(1 for _, _ok in _results if _ok)
print(f"  {_passed} of {len(_results)} checks passed")
print("=====")
if _passed == len(_results):
    print("Well done. Save the notebook and submit it.")
else:
    print("Read the FAIL lines above, fix those tasks, run this cell again.")

=====
SELF-CHECK   Practical 03 --- Choosing a Model Honestly
=====
[PASS] T1 | the depth-12 tree is trained
[PASS] T1 | T1_train_mae and T1_test_mae are that tree's two scores
[PASS] T1 | T1_gap is test minus train, and it is positive
[PASS] T2 | T2_scores has one score per depth
[PASS] T2 | every score is a plausible MAE in minutes
[PASS] T2 | T2_best_depth is the depth with the lowest score
[PASS] T3 | compare() returns one score per model
[PASS] T3 | the linear score matches the walkthrough
-----
8 of 8 checks passed
=====
Well done. Save the notebook and submit it.

```

Results

The four models on one split (80/20, seed 42)

Same 480 training orders, same 120 held-out orders, for every model.

Model	Train MAE	Test MAE	Gap
LinearRegression	2.04	1.92	-0.11
RandomForest (50 trees)	1.00	2.39	1.39
DecisionTree (no limit)	0.00	3.43	3.43
DecisionTree (depth 4)	3.66	4.23	0.57

The unlimited tree is overfitting in its purest form: a training MAE of **0.00** means it got all 480 training orders exactly right, and it still misses by 3.43 minutes on orders it has not seen. It memorised instead of learning.

The linear model's gap is **negative** (-0.11), meaning it scored slightly better on the test set than on the training set. That is not magic --- it just means the 120 test orders happened to be a touch easier, and the model is nowhere near memorising anything, so the two scores wobble around each other.

Cross-validation --- the same question, asked 5 times

Model	5-fold CV MAE	Spread (+/-)
LinearRegression	2.03	0.16
RandomForest (50 trees)	2.69	0.40
DecisionTree (depth 4)	4.57	0.18

Cross-validation moved linear regression from 1.92 to 2.03 and the forest from 2.39 to 2.69, so the single split had been flattering both of them a little. The ordering did not change, and the gap between linear and the forest (0.66) is much wider than either spread, so this ranking is real and not luck.

Task T1 --- overfitting measured as a gap

A `DecisionTreeRegressor(max_depth=12)`, scored by hand rather than with the helper:

MAE (minutes)	
Train	0.05
Test	3.40
Gap	3.35

Twelve levels of questions is already deep enough to almost memorise the training set (0.05 minutes of error, near-perfect), and the gap of 3.35 minutes is the price paid for it. It sits just under the unlimited tree's 3.43 --- capping the depth at 12 barely helped, because on 480 rows a depth-12 tree can still isolate nearly every order.

Task T2 --- tuning the forest with cross-validation only

`RandomForestRegressor(n_estimators=50)` at four depths, 5-fold CV, `X_test` and `y_test` never touched:

max_depth	CV MAE
2	5.271
4	3.643
6	2.844
None	2.688

Best depth: `None` --- the unlimited forest.

This is the opposite of what happened to the single tree in Step 10, where the best depth was 8 and removing the limit made things slightly worse. The reason is what a forest is for: 50 fully-grown trees each overfit in a *different* direction, and averaging them cancels most of that out. A forest does not need its depth capped the way a lone tree does.

Two honest caveats. The improvement from depth 6 to unlimited is only 0.156 minutes, which is well inside the forest's +/- 0.40 fold-to-fold spread --- so these two settings are effectively tied, not clearly ranked. And even at its best setting the forest's 2.69 still loses to plain linear regression's 2.03.

Task T3 --- the reusable comparison function

`compare({name: model})` fits each model on `X_train / y_train` and returns `{name: test MAE}`:

Model	Test MAE
linear	1.92
forest	2.39
tree4	4.23

The `linear` entry reproduces the walkthrough's 1.92 exactly, which is what the self-check verifies --- the function is doing the same work as Step 3, just for any number of models at once.

My choice

Model choice justification --- three sentences.

I would deploy `LinearRegression()` with scikit-learn's default settings, fitted on all four features (`distance_km` , `prep_time_min` , `traffic_level` , `rain`) and refitted on the full 600 orders once the choice is locked in. I am choosing it on the evidence that it has the lowest 5-fold cross-validated MAE of everything I tested --- **2.03 minutes (+/- 0.16)** against the cross-validated best forest, `RandomForestRegressor(n_estimators=50, max_depth=None)` , at **2.69 (+/- 0.40)**, a margin far wider than either fold-to-fold spread --- and, operationally, because it trains in milliseconds and ships as four coefficients rather than 50 stored trees, so it is cheap enough to retrain nightly, adds almost nothing to prediction latency, and can be explained to a restaurant owner in one line ("every extra kilometre adds about 3 minutes, rain adds about 5"). The trade-off I accept is that a straight line cannot represent any effect that is not linear: this dataset was generated from a linear formula, so it flatters my choice, and on messier real data with interactions --- rush-hour traffic behaving differently in the rain, say --- the forest would likely overtake it, at which point I would be knowingly trading accuracy away for the retraining cost and explainability above.

What to submit

1. This notebook, with every cell run and its output visible.
2. In a markdown cell at the end, three sentences: which model you would put in front of real customers, why, and what you gave up by choosing it.

Name your file `PXX_<your-roll-number>.ipynb` before you upload it.

How this practical is marked

What is marked	Marks
Walkthrough run end to end, outputs visible	3
Task T1 --- overfitting measured as a gap	2
Task T2 --- best setting found by cross-validation	3
Task T3 --- a reusable comparison function	2
Total	10

Read more

- scikit-learn --- Underfitting vs overfitting --- https://scikit-learn.org/stable/auto_examples/model_selection/plot_underfitting_overfitting.html
- scikit-learn --- Cross-validation --- https://scikit-learn.org/stable/modules/cross_validation.html
- scikit-learn --- Decision trees --- <https://scikit-learn.org/stable/modules/tree.html>
- scikit-learn --- Ensemble methods --- <https://scikit-learn.org/stable/modules/ensemble.html>
- matplotlib --- Quick start guide --- https://matplotlib.org/stable/users/explain/quick_start.html